

## CSE 2421/5042: Systems I

### Low-Level Programming and Computer Organization

# Introduction

## Presentation A

Gojko Babić

Read carefully: Bryant Chapter 1  
Study: Reek Chapter 2  
Skim: Reek Chapter 1

08/22/2018

## Some Basic Terms

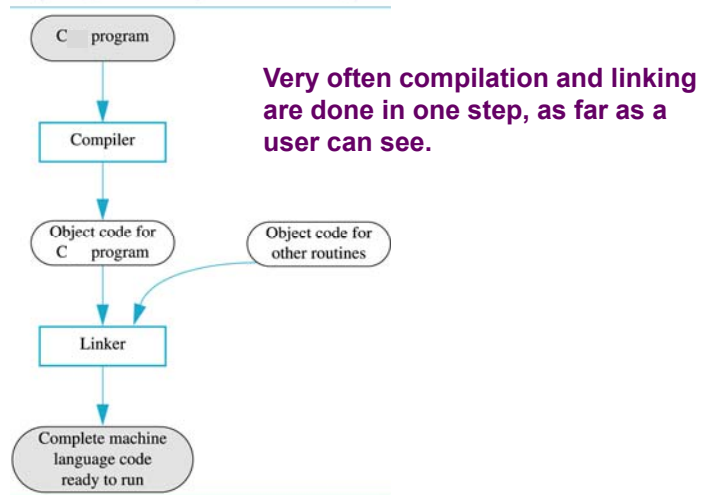
- **Algorithm**
  - a sequence of precise instructions which leads to a solution;
- **Program**
  - an algorithm expressed in a language the computer can understand;
- **Source code**
  - the original program in a high level language;
- **Compiler**
  - translates programs in high-level language to machine language;
- **Object code**
  - the (compiler) translated version in machine language;
- **Linker**
  - combines**
    - the object code for the programs we write,
  - and**
    - the object code for the pre-compiled routines,
  - into**
    - the executable code the CPU can run;

g. babic

Presentation A

2

## Preparing C Program for Running



g. babic

Presentation A

3

## Layout of a Simple C Program

```
/* include statements; if any
define statements; if any */

int main( )           // main function is mandatory
{
    Variable_Declarations

    Statement_1
    Statement_2
    ...
    Statement_Last

    return 0; // if you do not include this statement,
              // compiler will do it automatically
}
```

- The source code of a C program is stored in one or more files;
- Each function has to be completely contained in one file;
- Each source file is individually compiled to produce an object file;

g. babic

Presentation A

4

## Main Memory

- Long list of memory locations
  - each contains some pattern of zeroes and ones
  - can change during program execution
- Binary Digit or Bit: a digit that can only be zero or one
- Byte: each memory location has eight bits
- Address: a number that identifies a memory location
- Some data is too large for a single byte
  - e.g. most integers and real numbers are too large
  - next few consecutive bytes can store the additional bits for larger data
  - address of larger data refers to the first byte

g. babic

Presentation A

5

## Encoding Byte Values

- ❖ Byte = 8 bits
  - Binary (base 2 representation):  $00000000_2$  to  $11111111_2$ 
    - ✓ Use characters '0' and '1' → e.g. 01100101
  - Decimal (base 10 representation):  $0_{10}$  to  $255_{10}$ 
    - ✓ Use characters '0' to '9' → e.g. 127
  - Hexadecimal (base 16 representation):  $00_{16}$  to  $FF_{16}$ 
    - ✓ Use characters '0' to '9' and 'A' to 'F' → e.g. 0x7D
  - Octal (base 8 representation):  $000_8$  to  $377_8$ 
    - ✓ Use characters '0' to '7' → e.g. 075
  - Write  $FA1D37B_{16}$  in C → 0xFA1D37B or 0xfa1d37b → ? bits
  - Write  $114376_8$  in C → 0114376 → ? bits

Presentation A

6

## Unsigned Integer Encoding

- Here is how in C an 16-bit unsigned integer **variable** y can be declared and initialized to 15,213:  
`short unsigned int x = 15213;`
- But how is y represented as a binary number?

Let binary number X have w-bit pattern:  $x_{w-1} x_{w-2} x_{w-3} \dots x_3 x_2 x_1 x_0$   
 $x_i = 0 \text{ or } 1$

Binary to Unsigned decimal number interpretation:

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101

Presentation A

7

## Unsigned Integer Encoding Examples

**x =** 15213: 00111011 01101101

Weight	15213	
1	1	1
2	0	0
4	1	4
8	1	8
16	0	0
32	1	32
64	1	64
128	0	0
256	1	256
512	1	512
1024	0	0
2048	1	2048
4096	1	4096
8192	1	8192
16384	0	0
32768	0	0
<b>Sum</b>	<b>15213</b>	

8

## The hello program

```
#include <stdio.h>

int main()
{
    printf("hello, world\n");
    return 0;
}
```

ASCII text representation of hello.c program:

```
# i n c l u d e < s t d i o .
35 105 110 99 108 117 100 101 32 60 115 116 100 105 111 46
h > \n \n i n t < s p > m a i n ( ) \n {
104 62 10 10 105 110 116 32 109 97 105 110 40 41 10 123
\n < s p > < s p > < s p > < s p > p r i n t f ( " ...
10 32 32 32 32 112 114 105 110 116 102 40 34 ...
```

ASCII 7-bit code; ASCII is a part of UTF-8 that uses up to 4 bytes.

g. babic


Presentation A

9

## Compilation

- On Linux system, the translation from C program in the source file hello.c to executable object program in the file abc is performed by the following command:

```
unix> gcc -o abc hello.c
```



- To run executable file on a Unix system, type its name:

```
> abc
hello, world
>
```


g. babic

Presentation A

10

## More About gcc Command

- `-o file` → put the output in *file*  
`gcc -o abc hello.c`
- `-E` → stop after preprocessing; do not compile  
`gcc -E -o abc.i hello.c`
- `-S` → stop after compilation; do not assembly  
`gcc -S -o abc.s hello.c`
- `-c` → stop after assembly, do not link  
`gcc -c -o abc.o hello.c`



## hello.c x86-64 Assembly Code

```
#include <stdio.h>

int main()
{
    printf("hello, world\n");
    return 0;
}
```

```
main:
    subq $8, %rsp
    movl $.LC0, %edi
    call puts
    movl $0, %eax
    addq $8, %rsp
    ret
```

➤ This Linux gcc command:

`gcc -S -O1 hello.c`

produces x86-64 assembly code (64 bit code as default in our system) in file hello.s (on right) from C code in the file hello.c (on left) .

Note: switch `-O1` sets level 1 optimization

g. babic

12

## hello.c x86 Assembly Code

```
#include <stdio.h>
```

```
int main()
{
    printf("hello, world\n");
    return 0;
}
```

```
main:
```

```
    pushl %ebp
    movl %esp, %ebp
    andl $-16, %esp
    subl $16, %esp
    movl $.LC0, (%esp)
    call puts
    movl $0, %eax
    leave
    ret
```

- This Linux gcc command:  
`gcc -S -O1 -m32 hello.c`  
produces x86 assembly code (32 bit code)  
in file hello.s (on right) from C code in the  
file hello.c (on left).

Note: switch `-m32` produces 32 bit code

g. babic

13

## 64-bit Executable Code of hello.c Main Function

```
00000000004004c4 <main>:
```

4004c4:	48 83 ec 08	sub	\$0x8,%rsp
4004c8:	bf d8 05 40 00	mov	\$0x4005d8,%edi
4004cd:	e8 e6 fe ff ff	callq	4003b8 <puts@plt>
4004d2:	b8 00 00 00 00	mov	\$0x0,%eax
4004d7:	48 83 c4 08	add	\$0x8,%rsp
4004db:	c3	retq	
4004dc:	90	nop	
4004dd:	90	nop	
4004de:	90	nop	
4004df:	90	nop	

14

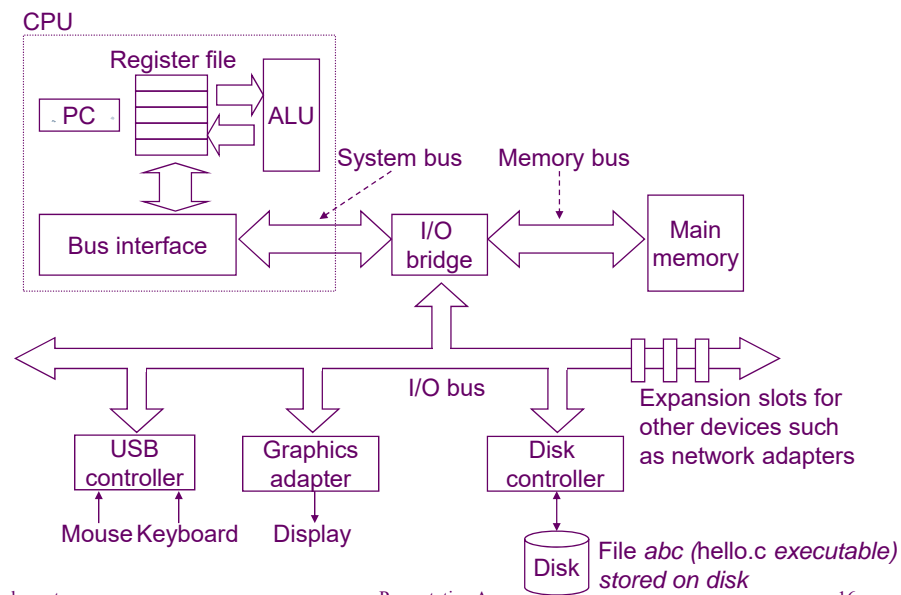
## 32-bit Executable Code of hello.c Main Function

```

080483b4 <main>:
80483b4:  55                push    %ebp
80483b5:  89 e5             mov     %esp,%ebp
80483b7:  83 e4 f0          and     $0xffffffff0,%esp
80483ba:  83 ec 10          sub     $0x10,%esp
80483bd:  c7 04 24 94 84 04 08 movl    $0x8048494,(%esp)
80483c4:  e8 27 ff ff ff    call   80482f0 <puts@plt>
80483c9:  b8 00 00 00 00    mov     $0x0,%eax
80483ce:  c9               leave
80483cf:  c3               ret
  
```

15

## Typical Computer Organization



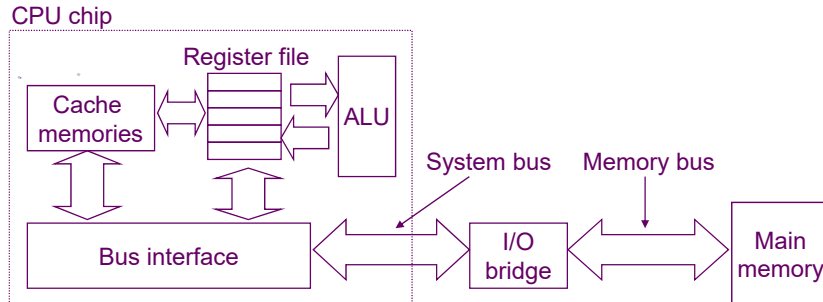
bryant

Presentation A

16



## Cache memory



Main memory size: in Giga bytes; 1GB = ??

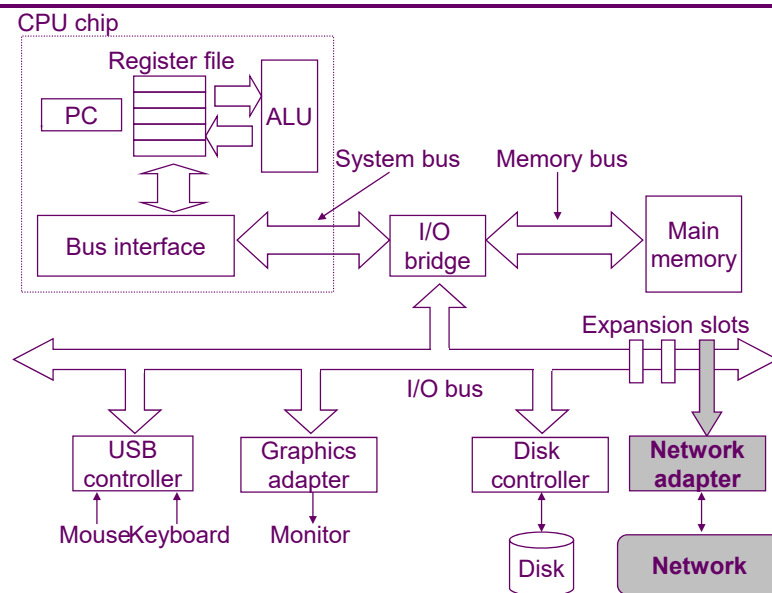
Cache size: in Mega bytes; 1MB = ??

Disk size: in Tera bytes; 1TB = ??

Presentation A

17

## Network Is Just Another I/O Device



Presentation A

18

## Operating System


- Protect the computer from misuse,
- Provide an abstraction for using the hardware so that programs can be written for a variety of different hardware,
- Operating system provide to user programs service of accessing I/O devices: keyboard, display, disc and network,
- Manage the resources to allow for reasonable use by all users and programs on a computer.
- Programs are often written as if they are the only things running on a system and OS allows them to work this way by providing an abstraction known as a process,
- Process is a running program (one or more threads of control), along with all the data associated with it ,
- OS uses context switching to give the appearance of multiple processes executing at once on a single processor.

g. babic

Presentation A

19

## Moving-Head Disk Mechanism



A sector size = 512B,

A sector (**only**) is a unit of transfer between memory and disc.

g. babic

Presentation A

20

## Reality No. 1: Computer Arithmetic

- ❖ Is  $x^2 \geq 0$ ?
  - Floating point numbers: Yes!
  - But, what are floating point numbers?
  - Integers:
    - $40000^2 \rightarrow 1600000000$
    - $50000^2 \rightarrow ??$
- ❖ Is  $(x + y) + z = x + (y + z)$ ?
  - Integers: Yes!
  - Floating point numbers:
    - $(1e20 - 1e20) + 3.14 \rightarrow 3.14$
    - $1e20 + (-1e20 + 3.14) \rightarrow ??$

bryant

Presentation A

21

## Reality No. 2: Assembly Programming

- ❖ Chances are, you'll never write programs in assembly  
Compilers are much better
- ❖ But: Understanding assembly is key to machine-level execution model
  - Behavior of programs in presence of bugs
    - ✓ high-level language models break down
  - Tuning program performance
    - ✓ understand optimizations done (or not done) by the compiler
    - ✓ understanding sources of program inefficiency
  - Implementing system software

bryant

Presentation A

22

## Reality No. 3: Memory Matters

---

- Memory is not unbounded
  - ✓ it must be allocated and managed
  - ✓ many applications are memory dominated
- Memory referencing bugs especially pernicious
  - ✓ effects are distant in both time and space
- Memory performance is not uniform
  - ✓ Cache, main memory and disk effects can greatly affect program performance
  - ✓ adapting program to characteristics of memory system can lead to major speed improvements

bryant

Presentation A

23

## Memory Reference Errors

---

- C (as well as and C++) doesn't provide memory protection:
  - ✓ out of bounds array references
  - ✓ invalid pointer values
  - ✓ abuses of C functions malloc/free
- Can lead to nasty bugs and whether or not bug has any effect depends on system and compiler
  - ✓ corrupted object logically unrelated to one being accessed
  - ✓ effect of bug may be first observed long after it is generated
- How can I deal with this?
  - ✓ understand what possible interactions may occur

bryant

Presentation A

24

## Reality No. 4: About Performance

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{ int i,j;
  for (i = 0; i < 2048; i++)
    for (j = 0; j < 2048; j++)
      dst[i][j] = src[i][j];
}
```

4.3ms

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{int i,j;
  for (j = 0; j < 2048; j++)
    for (i = 0; i < 2048; i++)
      dst[i][j] = src[i][j];
}
```

81.8ms

- Programs above run on 2.0 GHz Intel Core i7 Haswell

## The AdHello program

```
#include<stdio.h>
int main()
{
  char c1;
  printf("Enter 'a', 'b' or 'c':");
  scanf("%c", &c1);
  if (c1=='a') printf("hello A, world\n");
  else if (c1=='b') printf("hello B, world\n");
  else if (c1=='c') printf("hello C, world\n");
  else printf("Wrong input\n");
  return 0;
}
```

## Monday In Class Assignment

---

- Caldwell Lab room CL112
- Logon to Windows and then Linux.
- Create directory Cse2421 and subdirectory A1.
- In A1 directory, using one of Linux editors create hello.c; include your name at the beginning of the file.
- Compile hello.c and run it; you should get expected results.
- In the same directory, using one of Linux editors create AdHello.c; include your name at the beginning of the file.
- Compile AdHello.c and run it; you should get expected results.
- At the end of class, turn in hard copies of both source files.
- Also, electronically submit hello.c and AdHello.c using Carmen at Testing1. To activate Carmen on Linux server first issue command: %firefox & to run browser.

27